

# Transformations

Math 430, Winter 2017

# What do we do if our assumptions are violated?

1. Change your assumptions (hard, need more stats)
2. Transform  $y$ ,  $x$ , or both


# Transformations can...

- address non-linear patterns (i.e. linear on transformed scale)
- stabilize variance
- correct skew
- minimize the effects of outliers
- estimate percentage effects



Transforming the response

# Example

How is job difficulty related to salary?

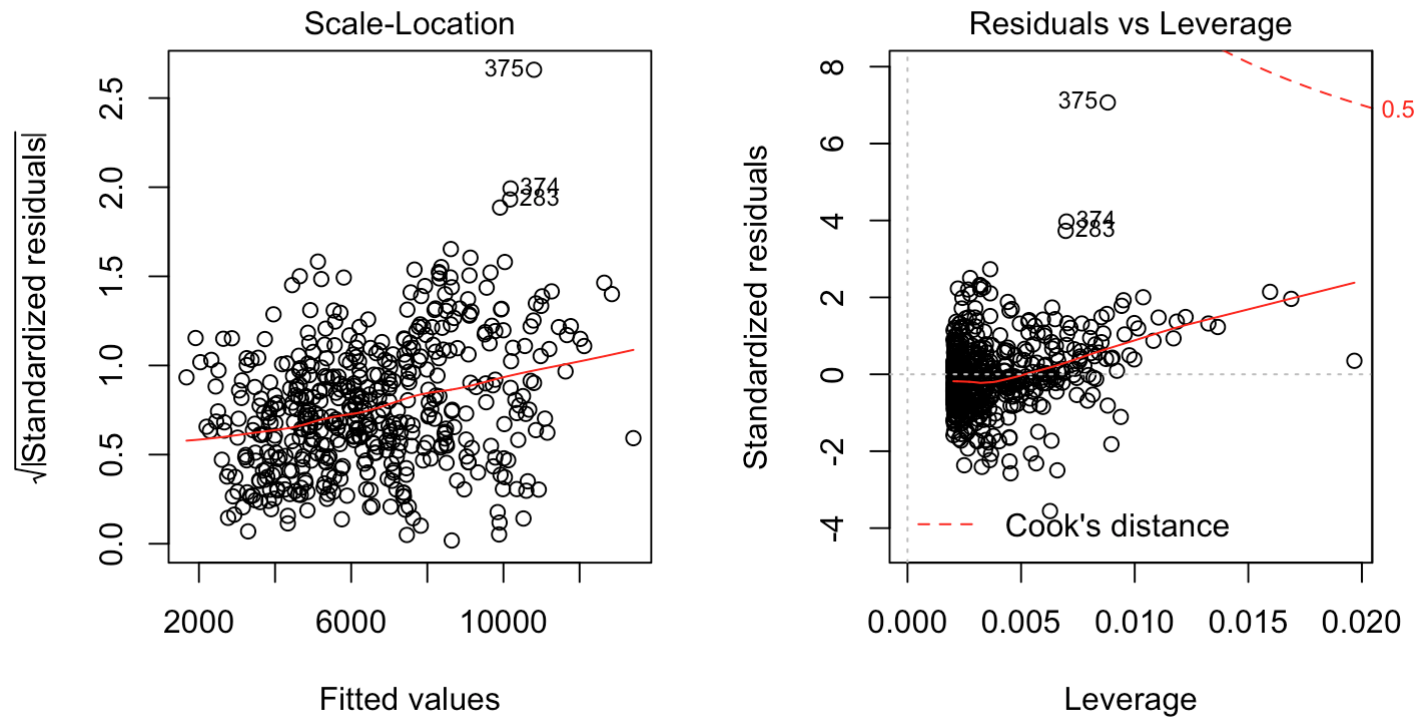


# Linearity and normality



Mean function may be nonlinear. Residuals have slightly heavy-tailed distribution.

# Constant variance and influence



Variance seems to be increasing. No troubling influential values.

# Power transformations

Let,  $U$  be a strictly positive variable, then the *power family of transformations* is defined by

$$\psi(U, \lambda) = U^\lambda$$

- Try values in the range  $[-1, 1]$  and see how they help problems,  $\{-1, -1/2, 0, 1/3, 1/2, 1\}$  is recommended
- Sometimes expansions to the range  $[-2, 2]$  is necessary



# Rules of thumb

- **log rule:** if values range over more than 1 order of magnitude and are strictly positive, then the natural log is likely helpful
- **range rule:** if the range is considerably less than 1 order of magnitude, then transformations are unlikely to help
- **square roots** are useful for count data

# Estimating $\lambda$

## Approach:


- Minimize  $RSS(\hat{\lambda})$
- Use an inverse response plot

## Application:

Use the `invResPlot` function found in the `car` package

# invResPlot

`invResPlot(salary_mod)` # prints lambda and RSS in console



# Transforming the linear model

## Approach 1: Create a new column in the data frame

```
library(dplyr)
salarygov <- mutate(salarygov, lmaxsalary = log(MaxSalary))
log_mod1 <- lm(lmaxsalary ~ Score, data = salarygov)
tidy(log_mod1)
```


##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	7.825242067	1.864010e-02	419.80678	0.000000e+00
## 2	Score	0.001889238	3.695849e-05	51.11783	3.553568e-199

## Approach 2: Apply the transformation in the `lm` formula


```
log_mod2 <- lm(log(MaxSalary) ~ Score, data = salarygov)
tidy(log_mod2)
```

##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	7.825242067	1.864010e-02	419.80678	0.000000e+00
## 2	Score	0.001889238	3.695849e-05	51.11783	3.553568e-199

# Transformed linear model?

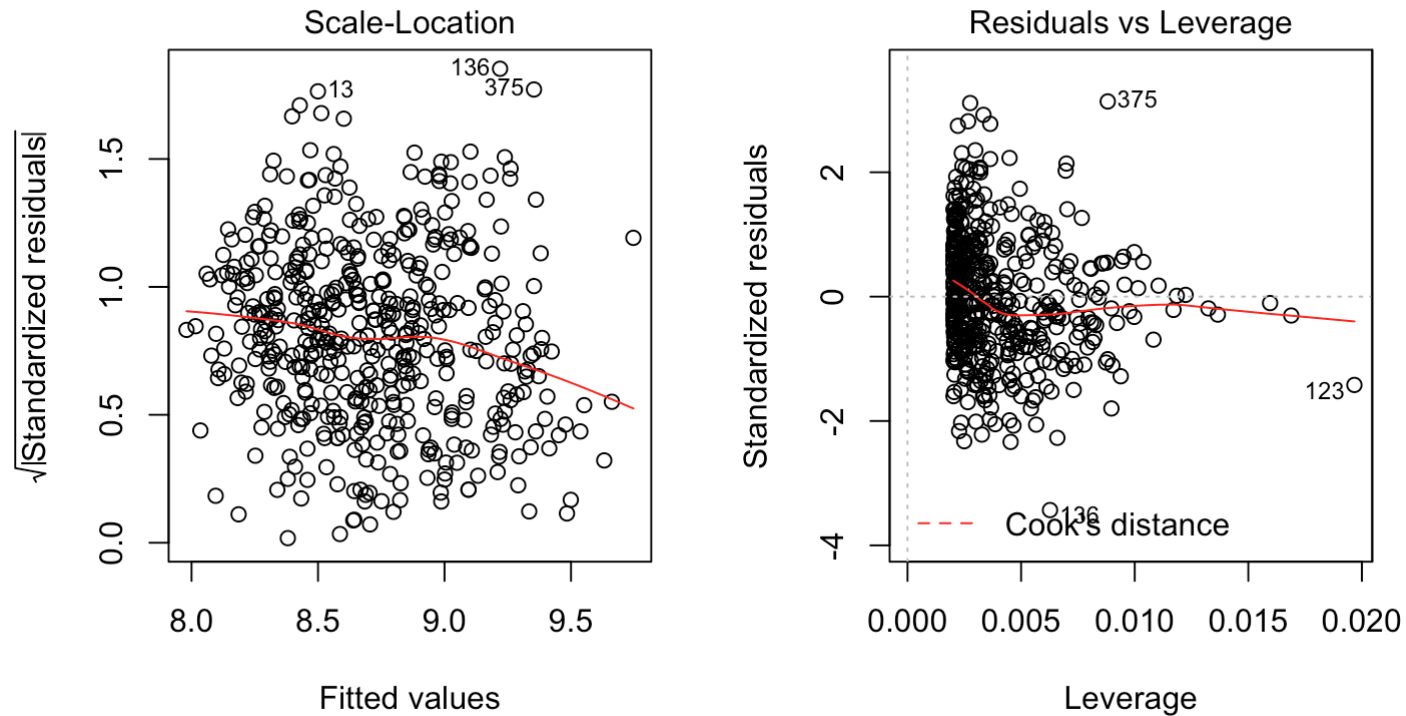


# Linearity and normality



The mean function appears to be linear and the residuals are well-approximated by the normal distribution.

# Constant variance and influence



There are no influential points and the appears to be constant.

# Interpreting the slope?



# Making predictions?

We can still use `predict` to make predictions...

```
newdata <- data.frame(Score = 469) # avg. difficulty score
cip <- predict(log_mod2, newdata = newdata, interval = "confidence")
cip
```

```
##          fit      lwr      upr
## 1 8.711295 8.697832 8.724757
```

but we need to back-transform

```
exp(cip)
```

```
##          fit      lwr      upr
## 1 6071.097 5989.913 6153.381
```

# Box-Cox transformation

We can automate the selection of the power transformation using the **modified power family** originally defined by Box and Cox (1964)


$$\psi_M(Y, \lambda_y) = \begin{cases} \text{gm}(Y)^{1-\lambda_y} \times (Y^{\lambda_y-1})/\lambda_y & \lambda_y \neq 0 \\ \text{gm}(Y) \times \log(Y) & \lambda_y = 0 \end{cases}$$

where  $\text{gm}(Y) = \exp(\sum \log(y_i)/n)$

- Transforming for normality of residuals
- $\lambda$  can be estimated via maximum likelihood

# Box-Cox transformation

`boxCox(salary_mod)` # to get the plot



# Box-Cox transformation

```
summary(powerTransform(salary_mod)) # print the estimate
```

```
## bcPower Transformation to Normality
```

```
##
```

```
##      Est.Power Std.Err. Wald Lower Bound Wald Upper Bound
```

```
## Y1      0.1292    0.0647              0.0025              0.256
```

```
##
```

```
## Likelihood ratio tests about transformation parameters
```

```
##              LRT df          pval
```


```
## LR test, lambda = (0)    3.927686  1 0.04749724
```

```
## LR test, lambda = (1) 179.872520  1 0.00000000
```

Transforming the predictor


# Example

How are tree height and tree diameter related for the Western red cedar?



# invTranPlot

`invTranPlot(Height ~ Dbh, data = ufcwc) # prints lambda and RSS in console`



# Transforming the linear model

## Approach 1: Create a new column in the data frame

```
ufcwc <- mutate(ufcwc, ldbh = log(Dbh))
rc_mod1 <- lm(Height ~ ldbh, data = ufcwc)
tidy(rc_mod1)
```

```
##           term  estimate std.error statistic    p.value
## 1 (Intercept) -463.3144  32.437870  -14.28313 6.505273e-29
## 2         ldbh   119.5192   5.531705   21.60621 5.761417e-46
```

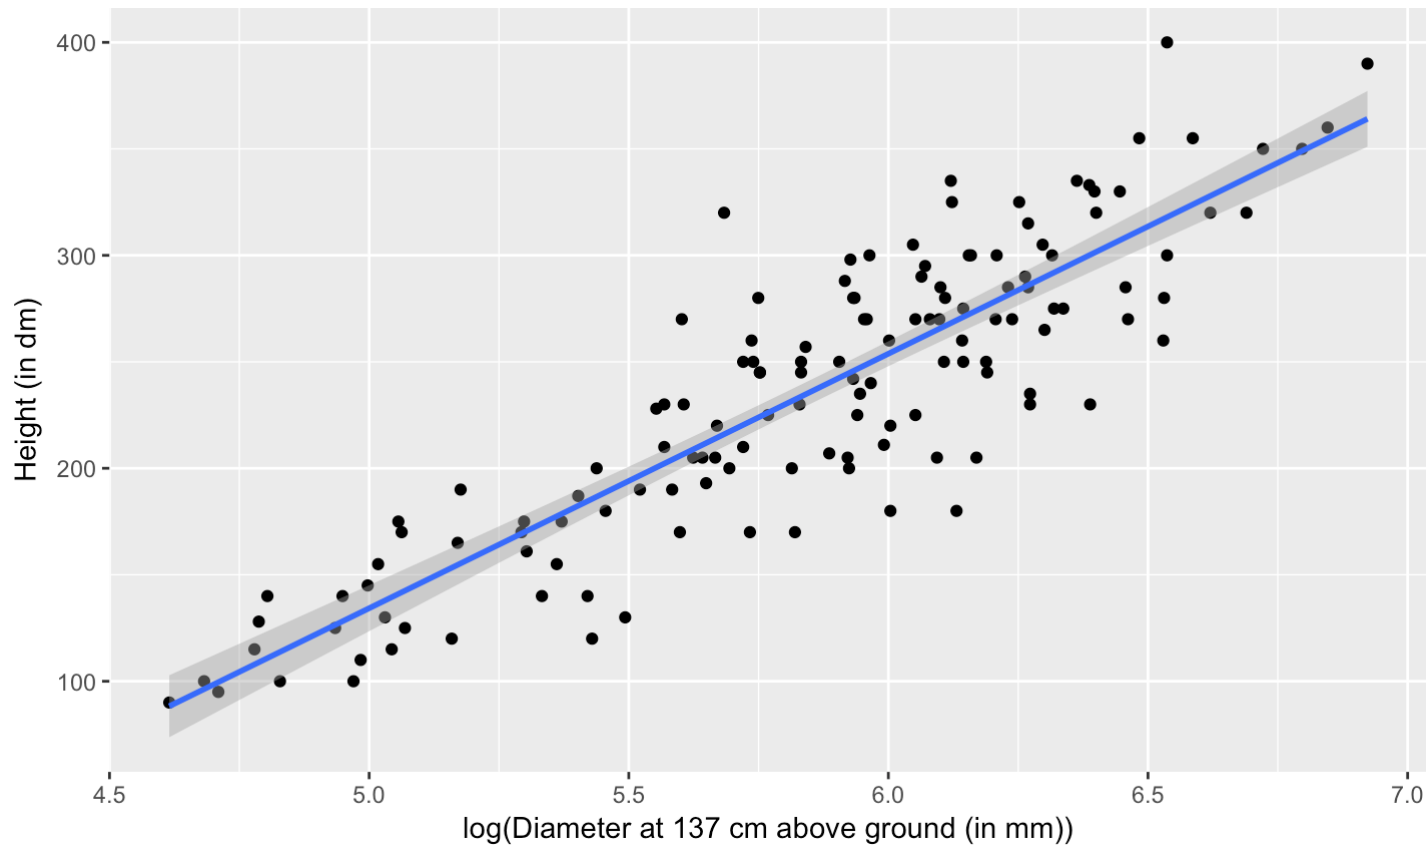
## Approach 2: Apply the transformation in the `lm` formula

```
rc_mod2 <- lm(Height ~ log(Dbh), data = ufcwc)
tidy(rc_mod2)
```


```
##           term  estimate std.error statistic    p.value
## 1 (Intercept) -463.3144  32.437870  -14.28313 6.505273e-29
## 2    log(Dbh)   119.5192   5.531705   21.60621 5.761417e-46
```



# Transformed linear model?

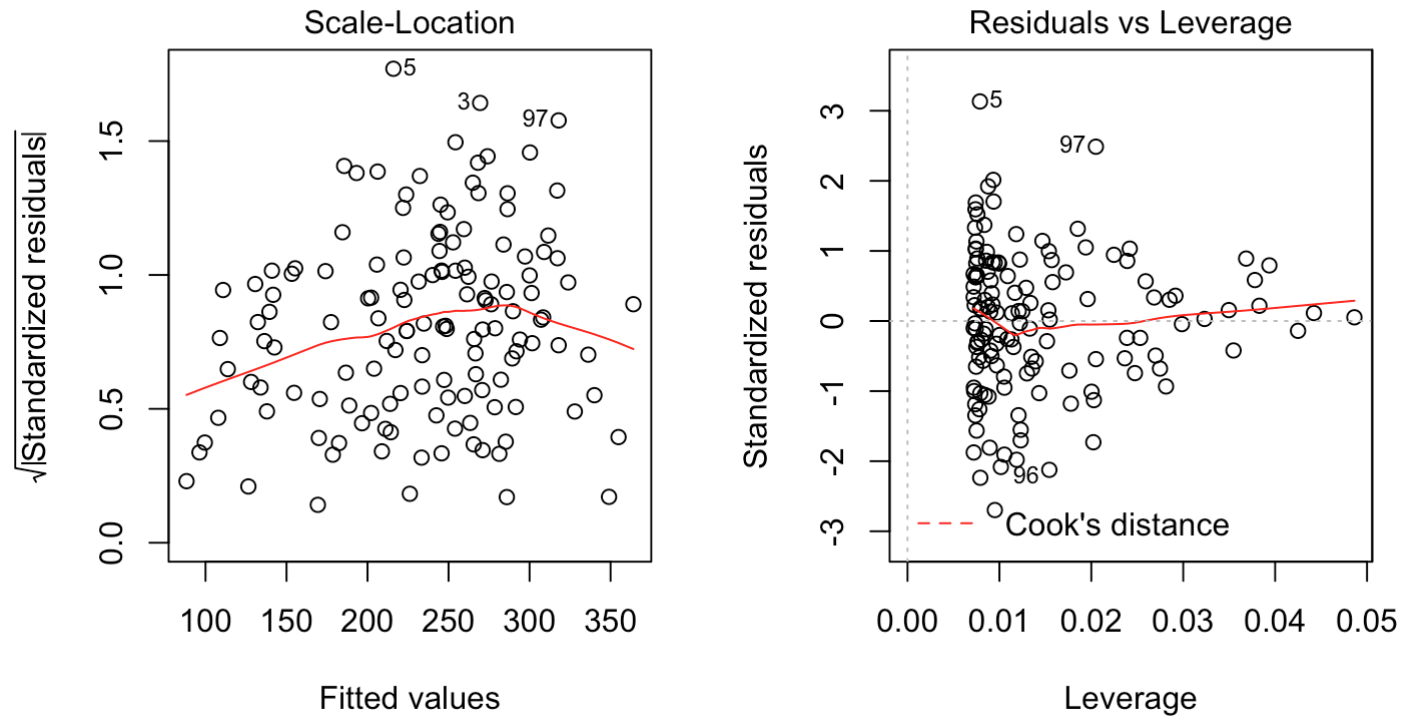


# Linearity and normality



The mean function appears to be linear and the residuals are well-approximated by the normal distribution.

# Constant variance and influence



There are no influential points and the appears to be roughly constant.

# Prediction

If you added a new column to the data frame...

```
newdata <- data.frame(ldbh = log(600))  
predict(rc_mod1, newdata = newdata, interval = "prediction")
```

```
##           fit          lwr          upr  
## 1 301.2415 234.8101 367.673
```

If you transformed `Dbh` in the model formula...


```
newdata2 <- data.frame(Dbh = 600)  
predict(rc_mod2, newdata = newdata2, interval = "prediction")
```

```
##           fit          lwr          upr  
## 1 301.2415 234.8101 367.673
```

Transforming both variables

# Example

How is the volume of a ship's cargo related to the time required to load and unload the cargo?



# Approaches

## Approach 1:

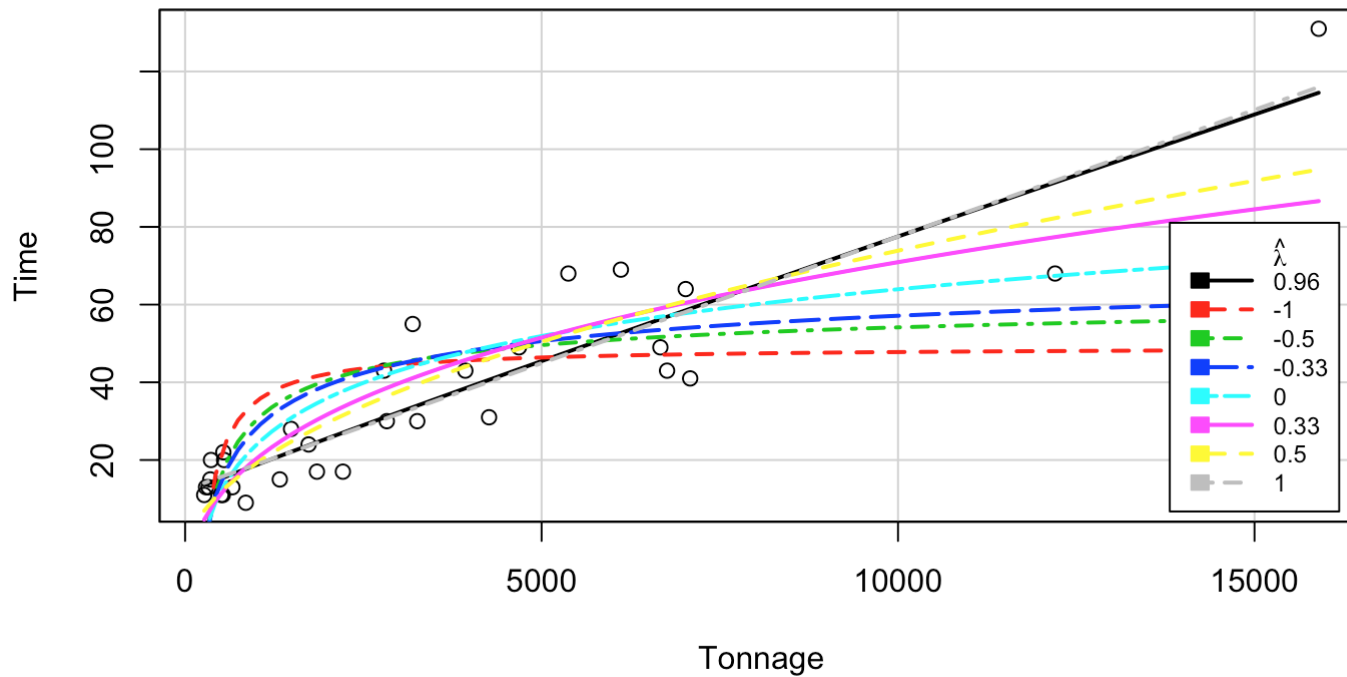
1. Use an inverse transformation plot to choose a transformation for X
2. Transform X, then use an inverse response plot to choose a transformation for Y

## Approach 2:

Transform X and Y simultaneously using the Box-Cox procedure

# Graphical approach

```
invTranPlot(Time ~ Tonnage, data = glakes, lambda = c(-1, -.5, -.33, 0, .33, .5, 1))
```




```
##      lambda      RSS
## 1  0.9591321  3313.093
## 2 -1.0000000 13096.852
```




# Graphical approach

```
cargo_mod1 <- lm(Time ~ I(Tonnage^.33), data = glakes)
invResPlot(cargo_mod1, lambda = c(-1, -.5, -.33, 0, .33, .5, 1))
```




# Transformed linear model?

```
cargo_mod2 <- lm(log(Time) ~ I(Tonnage.33), data = glakes)
```

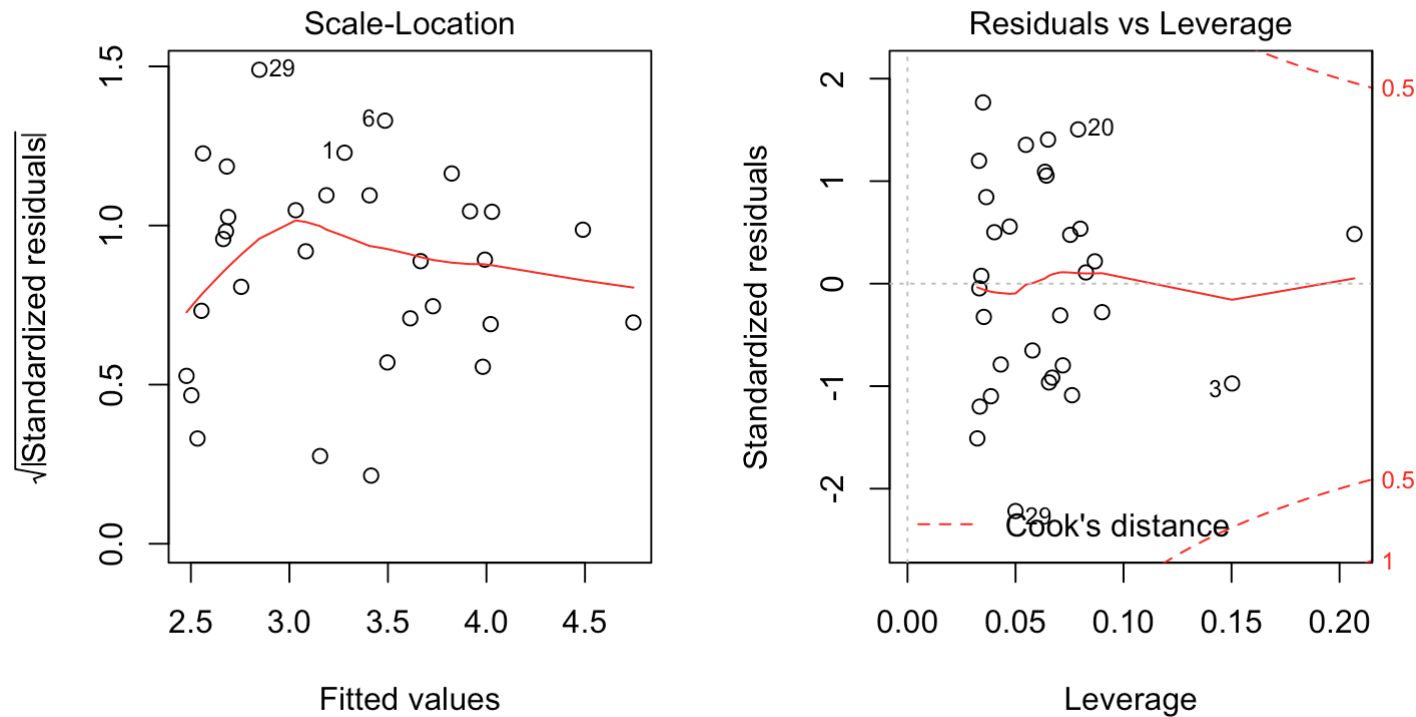


# Linearity and normality



The mean function appear to be linear and the residuals are well-approximated by the normal distribution.

# Constant variance and influence



There are no influential points and the appears to be roughly constant.

# Box-Cox approach

```
summary(powerTransform(cbind(Time, Tonnage) ~ 1, glakes))
```

```
## bcPower Transformations to Multinormality
```

```
##
```

```
##           Est.Power Std.Err. Wald Lower Bound Wald Upper Bound
```

```
## Time           0.0228   0.1930           -0.3554           0.4011
```

```
## Tonnage         0.2378   0.1237           -0.0046           0.4802
```

```
##
```

```
## Likelihood ratio tests about transformation parameters
```


```
##                LRT df                pval
```

```
## LR test, lambda = (0 0)  3.759605  2 1.526202e-01
```


```
## LR test, lambda = (1 1) 45.315290  2 1.445140e-10
```

# Transformed linear model?

```
cargo_bcmmod <- lm(log(Time) ~ I(Tonnage.25), data = glakes)
```

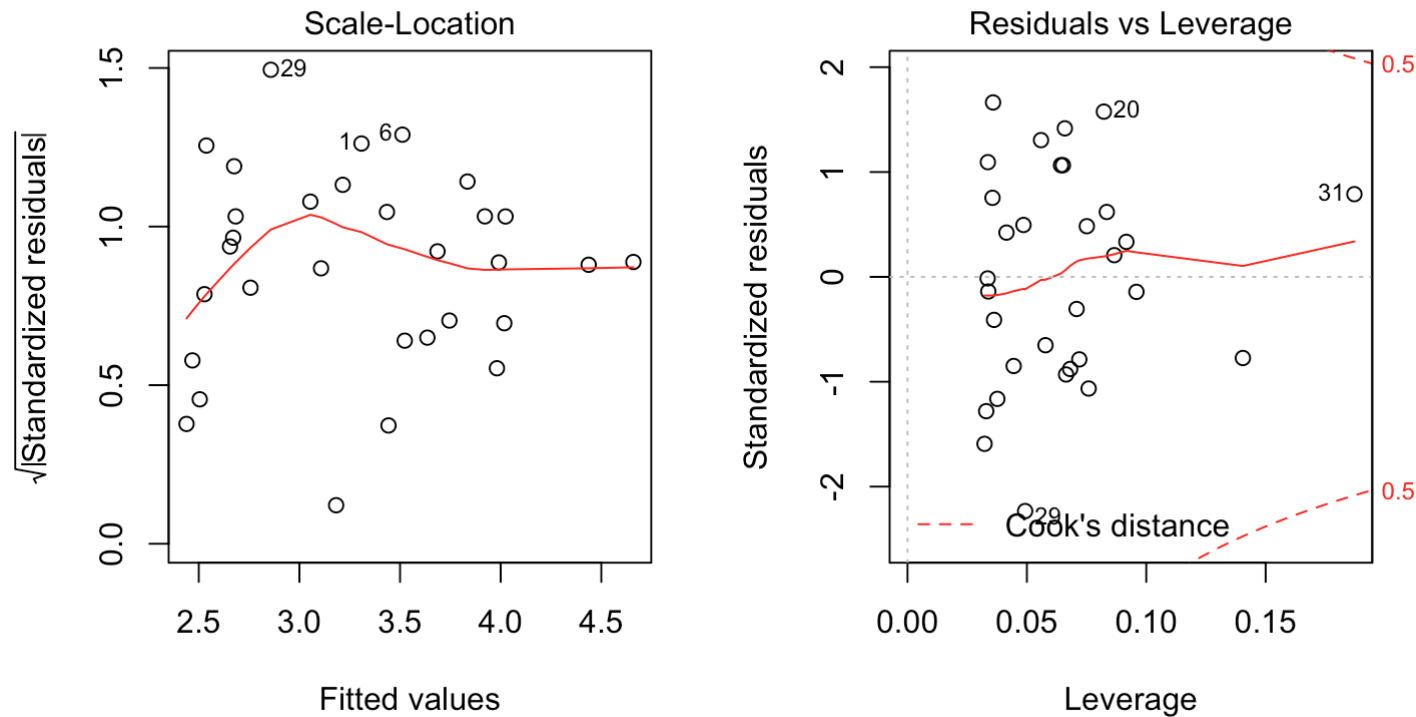


# Linearity and normality



The mean function appears to be linear and the residuals are well-approximated by the normal distribution.

# Constant variance and influence



There are no influential points and the appears to be roughly constant.



# A note on log transformations

If we apply a log transformation to both the response and the predictor, then

$$\% \Delta Y \approx \beta_1 \times \% \Delta x$$

- So, for every 1% increase in  $x$ , the model predicts a  $\beta_1$  % increase in  $Y$
- $\beta_1$  needs to be small for this to work out (see p.79 for details)

# Issues with Transformations

- You're often guessing
  - Statistics is an art AND a science!
- Changes the interpretation of the parameters
  - need to back-transform to provide interpretable results
- Changes SEs of the parameters
- Not always easy to keep track of all your assumptions