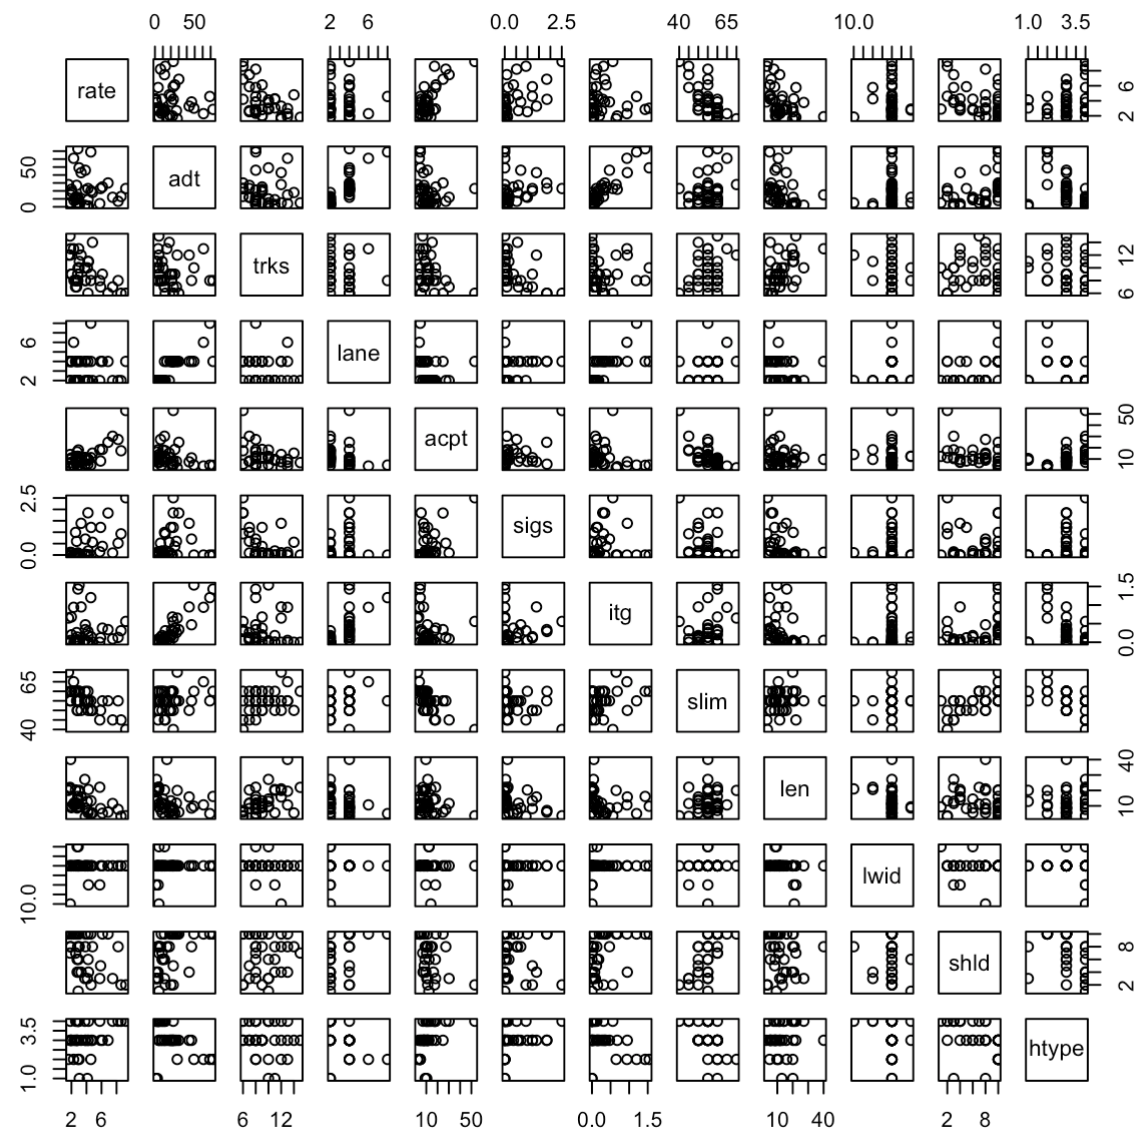



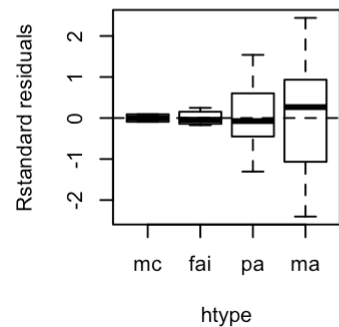

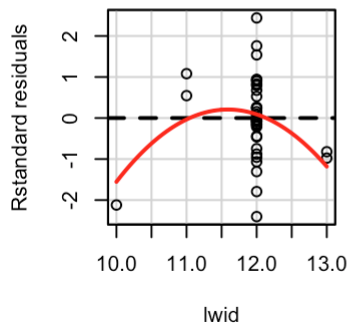




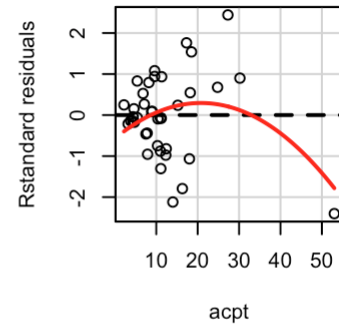



# MLR: Model Selection in R

Math 430, Winter 2017

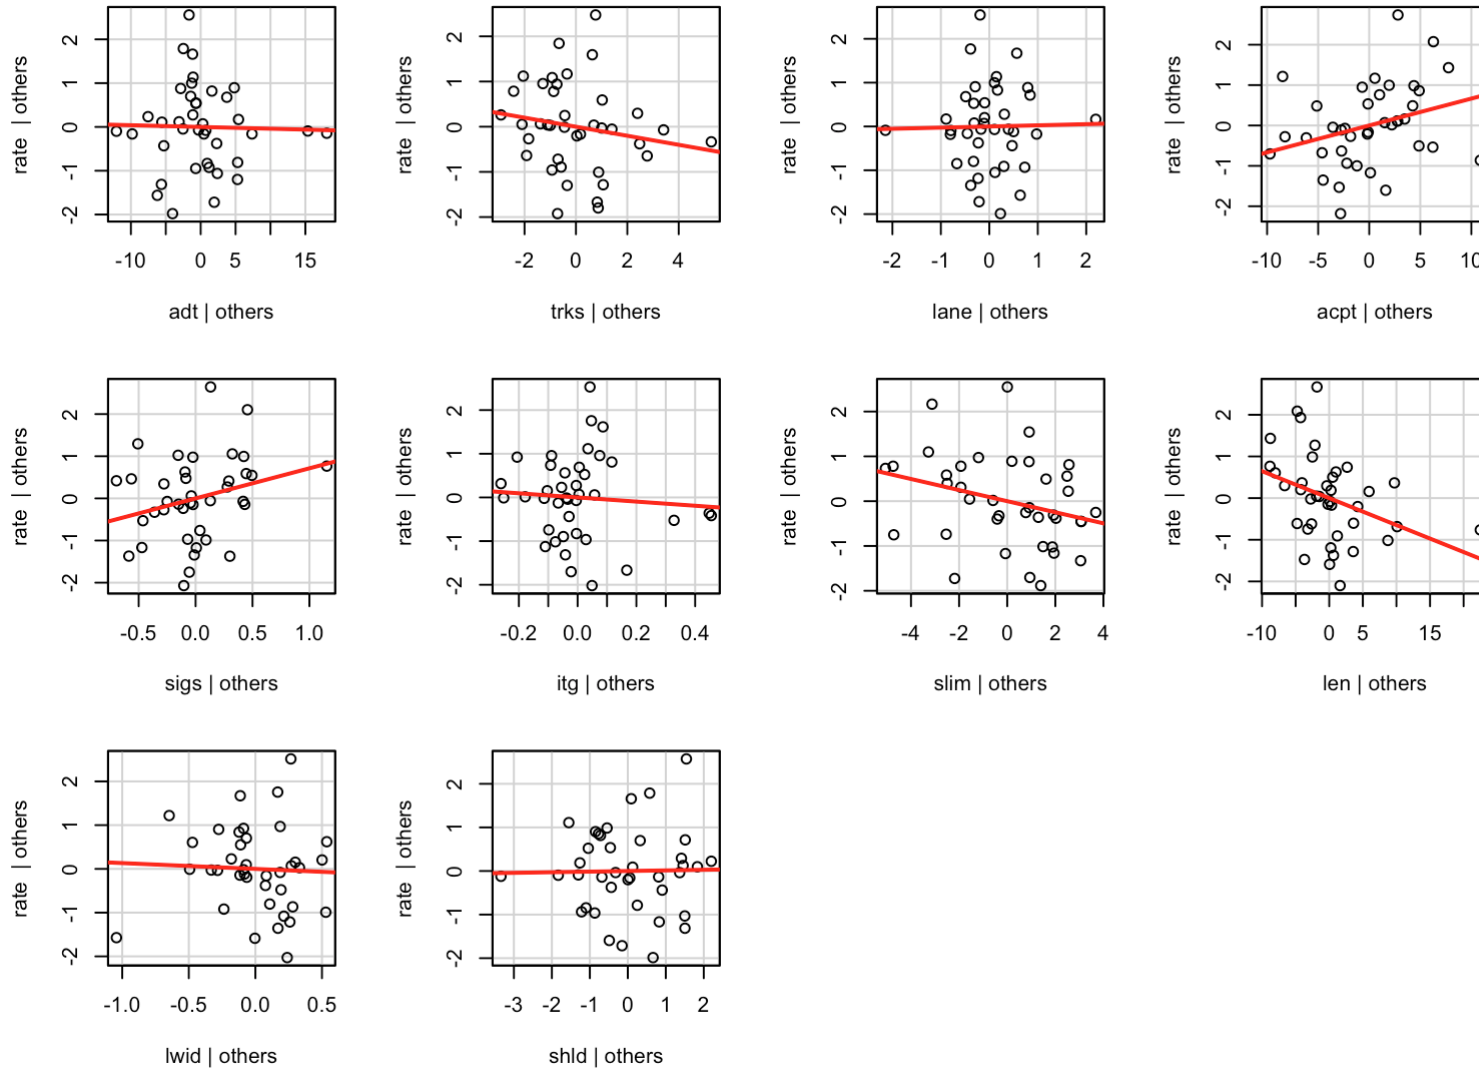
# Highway accident data

Variable	Description
<b>adt</b>	average daily traffic count (thousands)
<b>trks</b>	truck volume as a percent of the total volume
<b>lane</b>	total number of lanes of traffic
<b>acpt</b>	number of access points per mile
<b>sigs</b>	number of signalized interchanges per mile
<b>itg</b>	number of freeway-type interchanges per mile
<b>slim</b>	speed limit
<b>len</b>	length of the Highway segment (miles)
<b>lwid</b>	lane width (feet)
<b>shld</b>	width in feet of outer shoulder on the roadway
<b>htype</b>	type of roadway/funding source
<b>rate</b>	accident rate per million vehicle miles

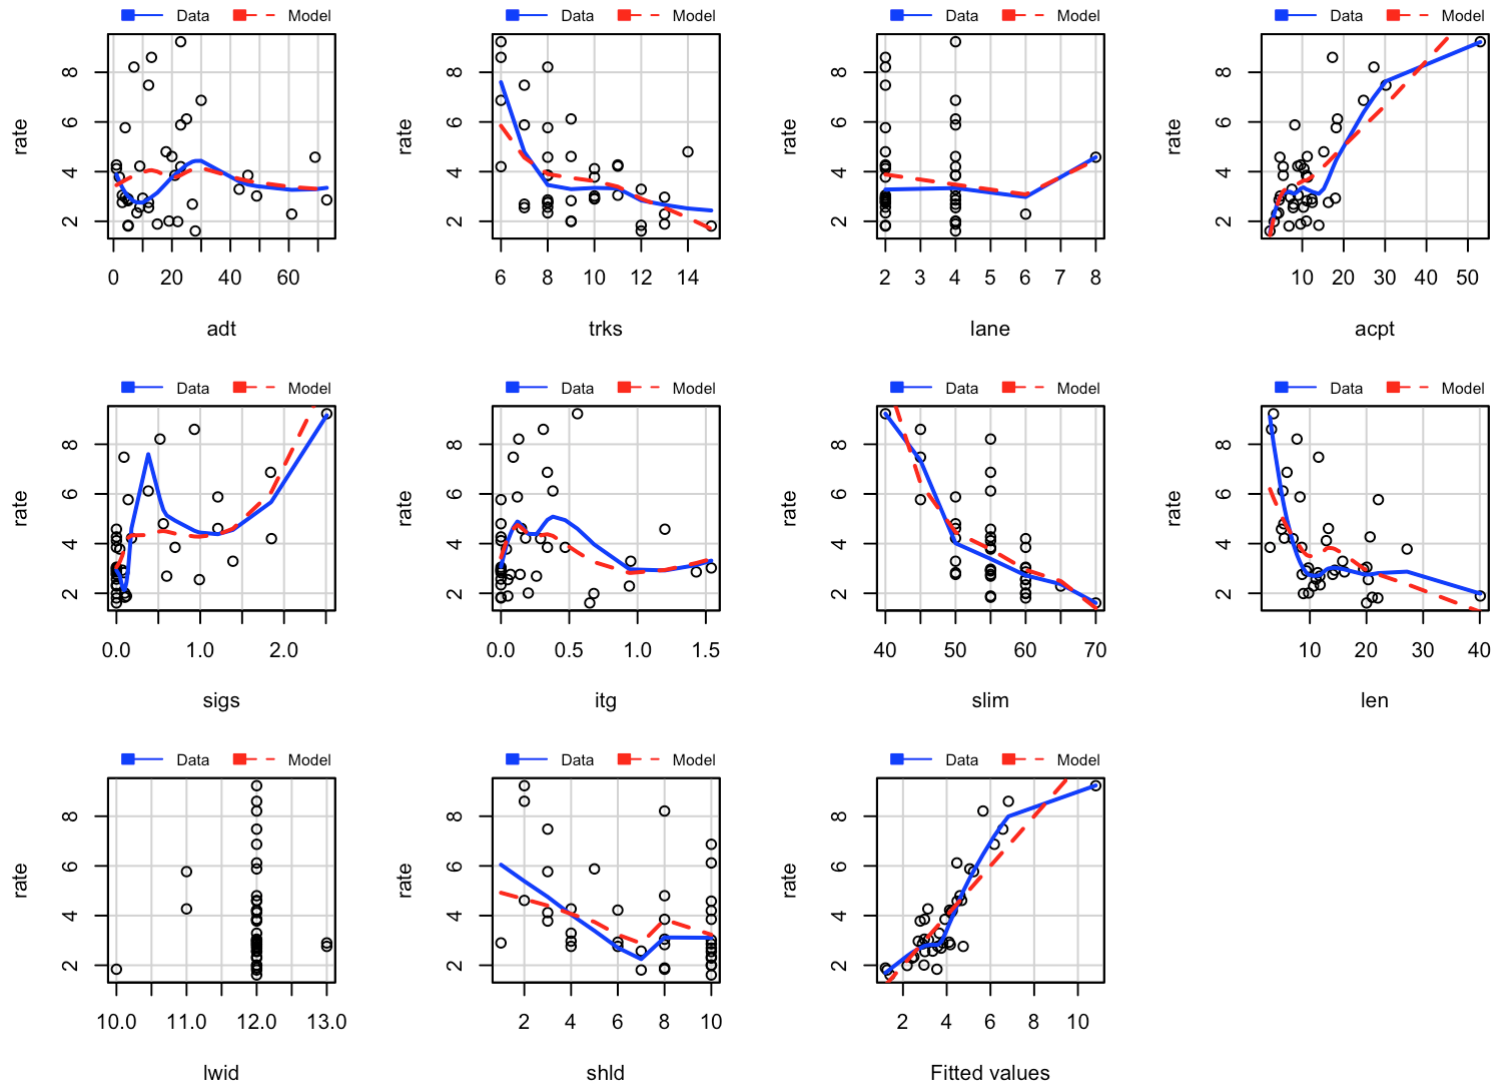


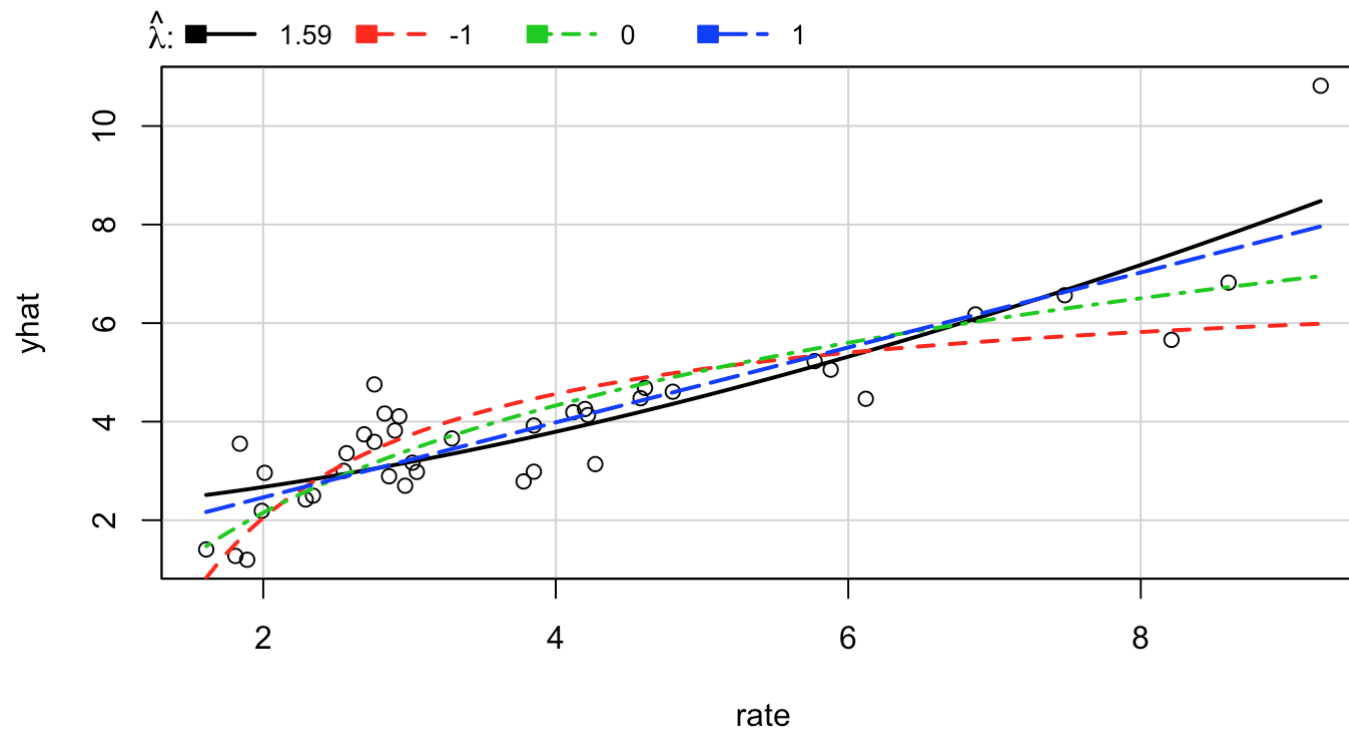


## Added-Variable Plots



## Marginal Model Plots


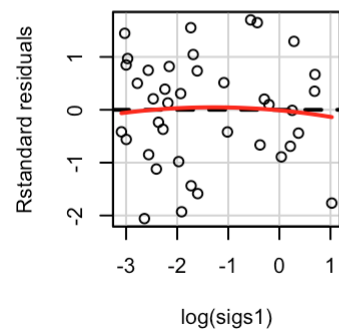

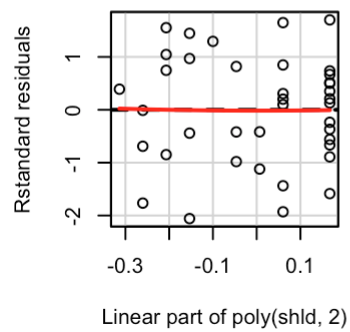
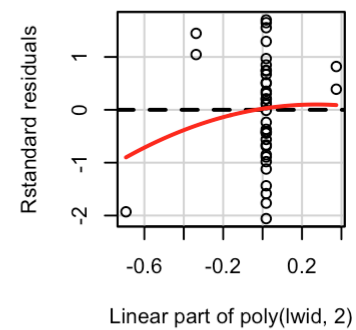











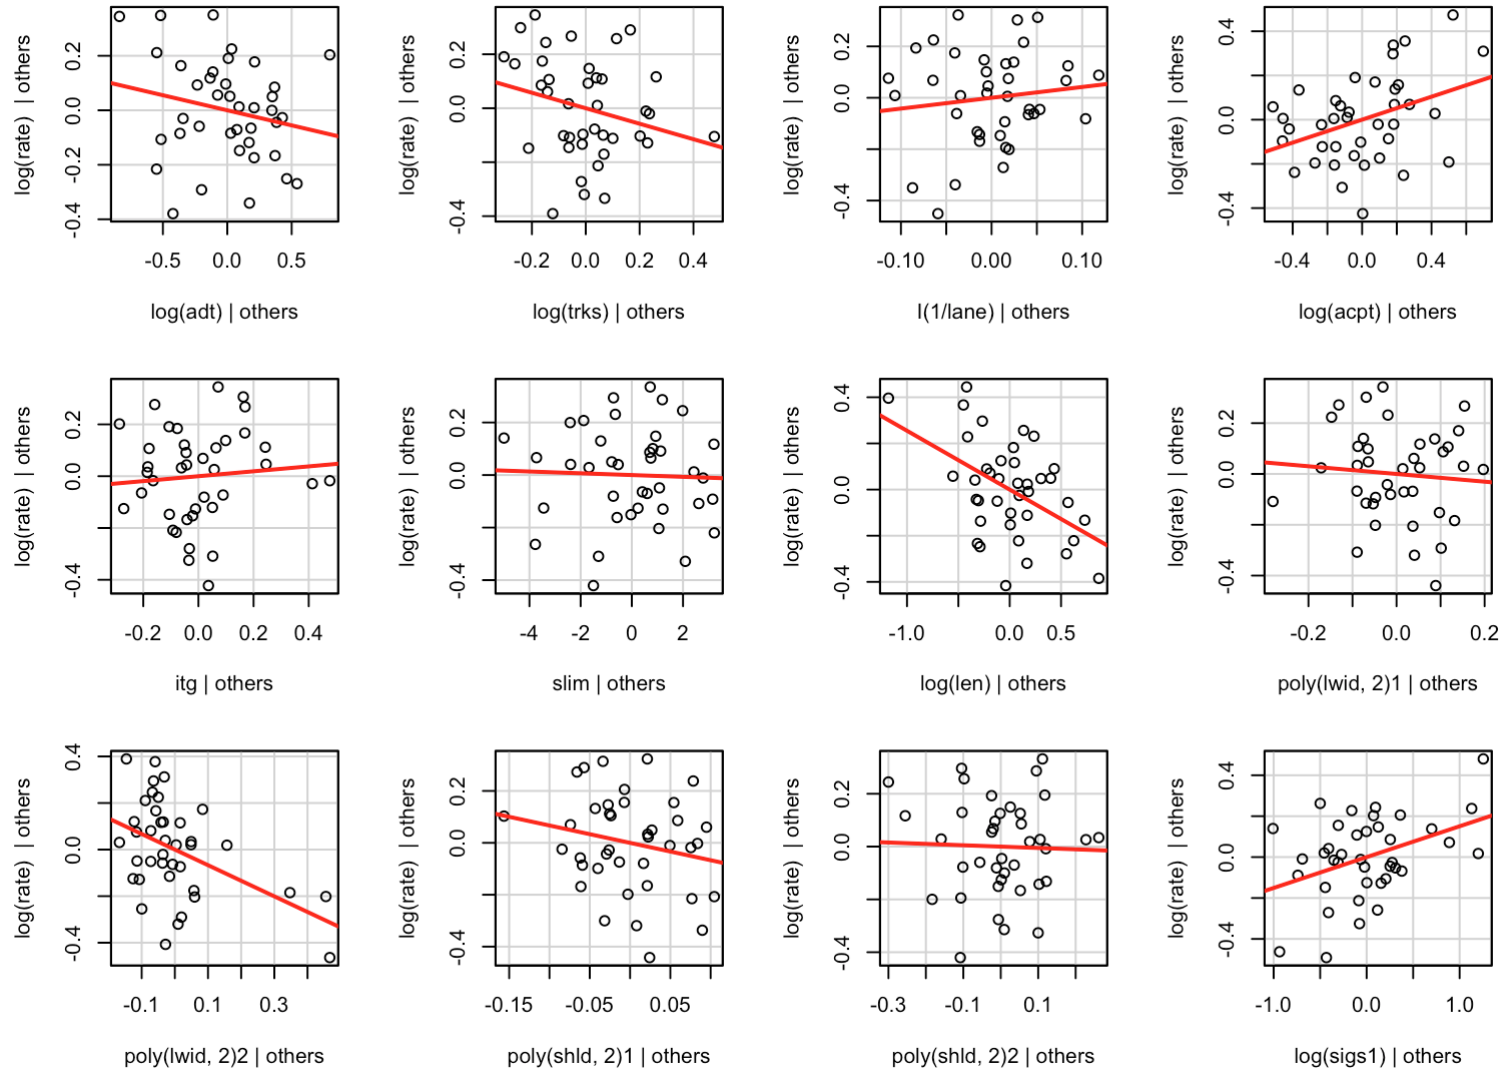
##	lambda	RSS
## 1	1.594529	26.41210
## 2	-1.000000	45.59093
## 3	0.000000	33.75636
## 4	1.000000	27.29810



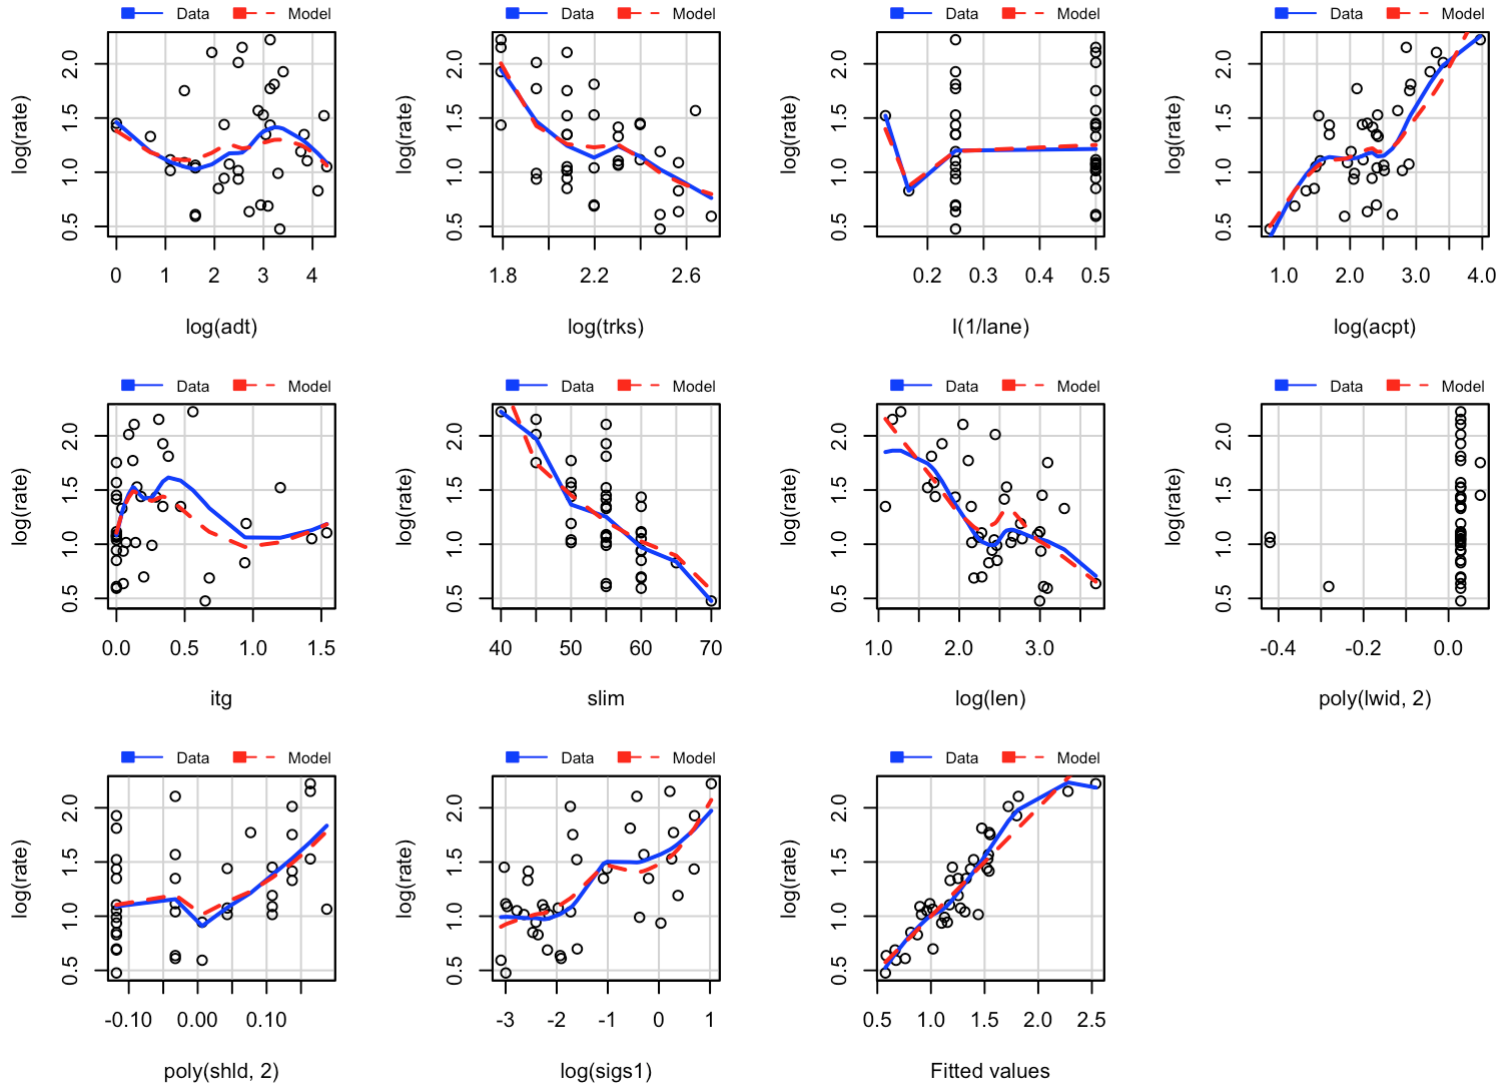
```
Highway <- mutate(Highway, sigs1 = (sigs * len + 1)/len)
full_mod_tform <- lm(log(rate) ~ log(adtt) + log(trks) + I(1/lane) + log(acpt) +
  itg + slim + log(len) + poly(lwid, 2) + poly(shld, 2) + htype + log(sigs1),
  data = Highway)
```



## Added-Variable Plots



## Marginal Model Plots



The **step** command

# Backward elimination

```
belim <- step(full_mod_tform, scope = list(lower = ~ 1), direction = "backward")
```

```
broom::tidy(belim)
```

##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	3.24639448	0.75119775	4.3216244	0.0001764473
## 2	log(adl)	-0.14429407	0.07746273	-1.8627547	0.0730195963
## 3	log(acpt)	0.18987179	0.10707212	1.7733075	0.0870567742
## 4	slim	-0.02011261	0.01007683	-1.9959263	0.0557516497
## 5	log(len)	-0.25644916	0.07871784	-3.2578279	0.0029403083
## 6	poly(lwid, 2)1	0.13688282	0.25106602	0.5452065	0.5899285279
## 7	poly(lwid, 2)2	-0.60177023	0.23510121	-2.5596220	0.0161662281
## 8	htypefai	0.33059140	0.33000676	1.0017716	0.3250331856
## 9	htypepa	-0.21786065	0.21955592	-0.9922786	0.3295598277
## 10	htypema	-0.06105924	0.18951707	-0.3221833	0.7497070874
## 11	log(sigs1)	0.17789568	0.05689946	3.1264916	0.0040983118

# Forward selection

```
null_mod <- lm(slog(rate) ~ 1, data = Highway)
fselect <- step(null_mod, scope = list(lower = ~ 1,
upper = ~ log(adl) + log(trks) + lane + acpt + itg + slim + log(len) +
      lwid + shld + htype + sigsl),
direction = "forward")
```

```
broom::tidy(fselect)
```

##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	4.16654113	0.741064508	5.622373	2.666474e-06
## 2	slim	-0.03185201	0.010261763	-3.103951	3.832823e-03
## 3	log(len)	-0.23573454	0.084896763	-2.776720	8.867186e-03
## 4	acpt	0.01100449	0.006669289	1.650025	1.081474e-01
## 5	log(trks)	-0.32903691	0.213483661	-1.541274	1.325068e-01

# Stepwise selection

```
step_hwy <- step(null_mod, scope = list(lower = ~ 1,  
upper = ~ log(adtl) + log(trks) + lane + acpt + itg + slim + log(len) +  
      lwid + shld + htype + sigsl),  
direction = "both")
```

```
broom::tidy(step_hwy)
```

##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	4.16654113	0.741064508	5.622373	2.666474e-06
## 2	slim	-0.03185201	0.010261763	-3.103951	3.832823e-03
## 3	log(len)	-0.23573454	0.084896763	-2.776720	8.867186e-03
## 4	acpt	0.01100449	0.006669289	1.650025	1.081474e-01
## 5	log(trks)	-0.32903691	0.213483661	-1.541274	1.325068e-01



# Using BIC rather than AIC

```
belim_bic <- step(full_mod_tform, scope = list(lower = ~ 1), direction = "backward",  
               k = log(nrow(Highway)))
```

```
broom::tidy(belim_bic)
```

##	term	estimate	std.error	statistic	p.value
## 1	(Intercept)	3.24639448	0.75119775	4.3216244	0.0001764473
## 2	log(adl)	-0.14429407	0.07746273	-1.8627547	0.0730195963
## 3	log(acpt)	0.18987179	0.10707212	1.7733075	0.0870567742
## 4	slim	-0.02011261	0.01007683	-1.9959263	0.0557516497
## 5	log(len)	-0.25644916	0.07871784	-3.2578279	0.0029403083
## 6	poly(lwid, 2)1	0.13688282	0.25106602	0.5452065	0.5899285279
## 7	poly(lwid, 2)2	-0.60177023	0.23510121	-2.5596220	0.0161662281
## 8	htypefai	0.33059140	0.33000676	1.0017716	0.3250331856
## 9	htypepa	-0.21786065	0.21955592	-0.9922786	0.3295598277
## 10	htypema	-0.06105924	0.18951707	-0.3221833	0.7497070874
## 11	log(sigs1)	0.17789568	0.05689946	3.1264916	0.0040983118

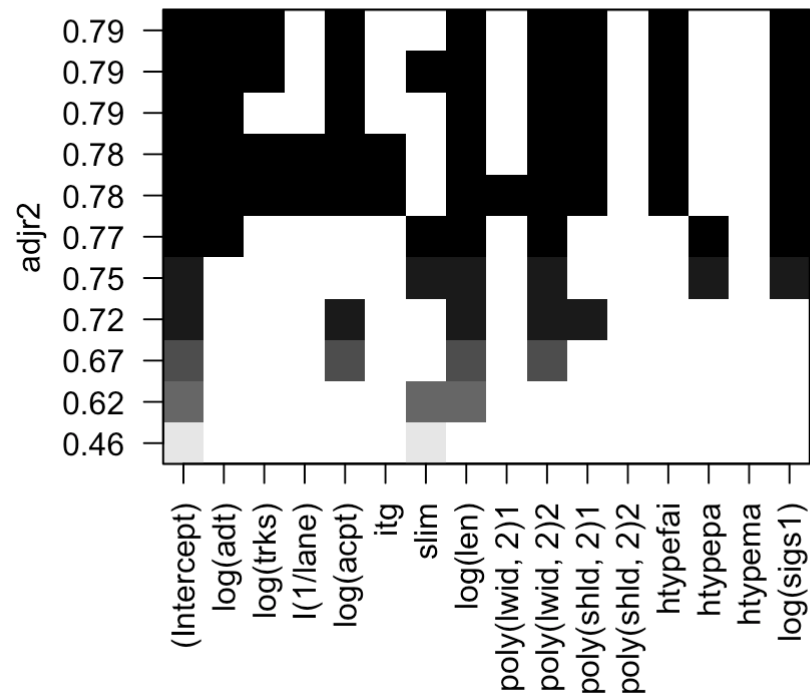
The **regsubsets** command

# All subsets in R

```
library(leaps)
regfit_full <- regsubsets(log(rate) ~ log(adtl) + log(trks) + I(1/lanel) + log(acpt) +
  itg + slim + log(len) + poly(lwid, 2) + poly(shld, 2) + htype + log(sigs1),
  data = Highway, method = "exhaustive", nvmax = 11, nbest = 1)
reg_summary <- summary(regfit_full)
```


# Investigating the results

```
plot(regfit_full, scale = "adjr2")
```

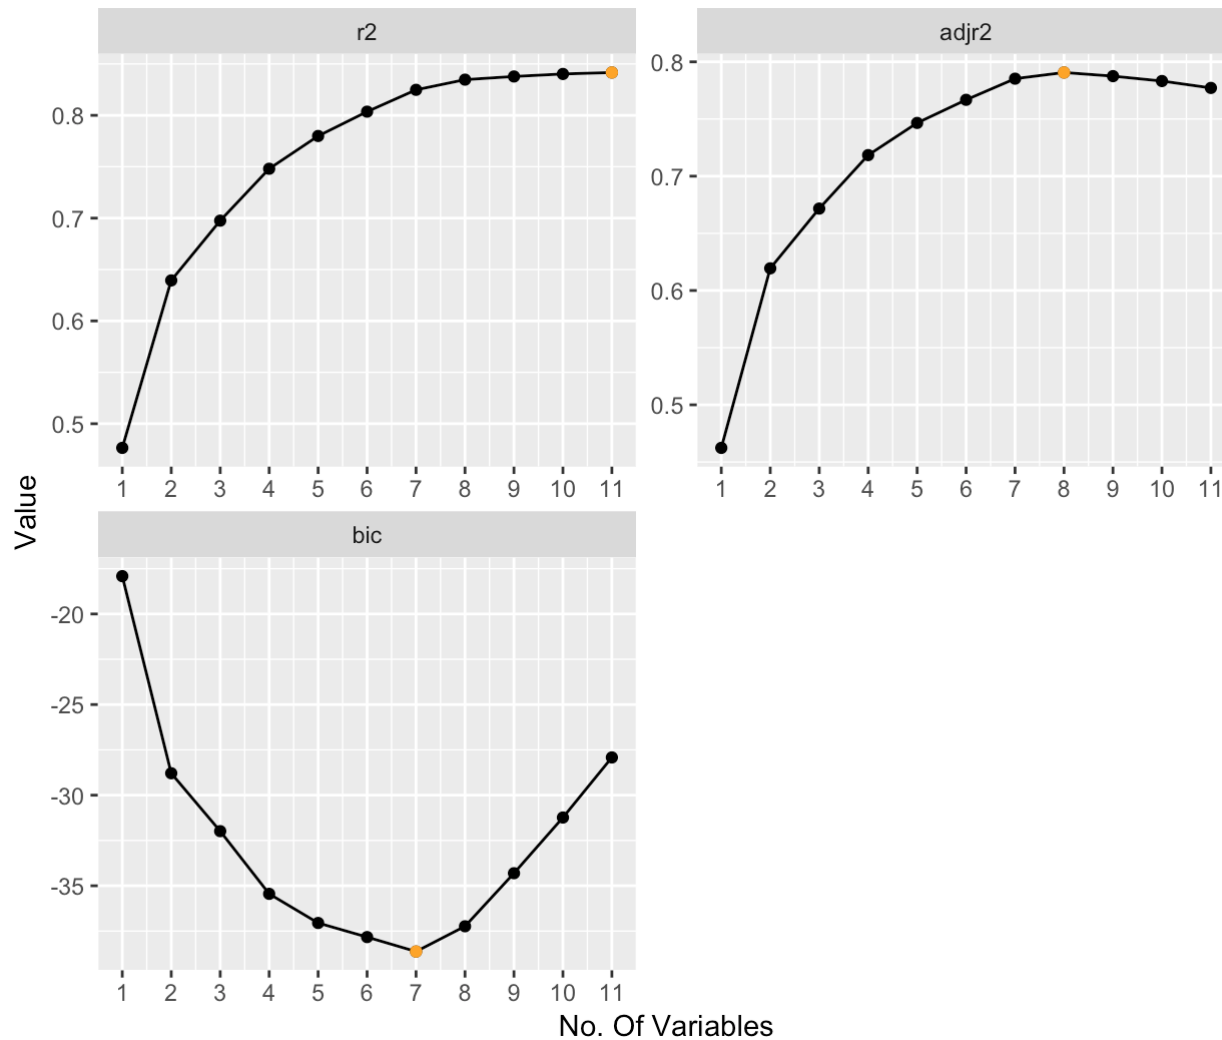


# Investigating the results

```
plot(regfit_full, scale = "bic")
```



# Another plot option



Extracting goodness-of-fit  
measures

```
broom::glance(step_hwy)
```

```
##   r.squared adj.r.squared      sigma statistic      p.value df    logLik
## 1 0.6960742      0.6603182 0.2698162   19.46735 2.06722e-08   5 -1.572629
##           AIC          BIC deviance df.residual
## 1 15.14526 25.12663 2.475227           34
```

Extract  $R^2_{adj}$

```
broom::glance(step_hwy)$adj.r.squared
```

```
## [1] 0.6603182
```



## Calculate AIC

```
# The first number is equiv. d.f., the second is AIC  
extractAIC(step_hwy, k = 2)
```

```
## [1] 5.00000 -97.53195
```

## Calculate AICc

```
n <- nrow(Highway)  
nslope <- length(step_hwy$coefficients) - 1  
extractAIC(step_hwy, k = 2) + 2 * (nslope + 1) * (nslope + 2) / (n - nslope - 1)
```

```
## [1] 6.764706 -95.767241
```

## Calculate BIC

```
extractAIC(step_hwy, k = log(n))
```

```
## [1] 5.00000 -89.21414
```

# Training and test data sets

*# Select rows for a training data set*

```
train_id <- sample(1:nrow(df), size = round((2/3) * nrow(df)))
```

*# Create the training and test data sets*

```
train <- df[train_id,]
```

```
test  <- df[-train_id,]
```